

EE359 Big Data Mining

# Streaming Algorithms

Jiaxin Ding

John Hopcroft Center



上海交通大学  
约翰·霍普克罗夫特  
计算机科学中心

John Hopcroft Center for Computer Science



# Data Streams

- A data stream is a **sequence** of signals used to transmit or receive information that is in the process of being **transmitted**. In many situations, we do not know the entire data set in advance.
  - **Infinite**
  - **Non-stationary**

. . . 1, 5, 2, 7, 0, 9, 3

. . . a, r, v, t, y, h, b

. . . 0, 0, 1, 0, 1, 1, 0  
**time**

# The Stream Model

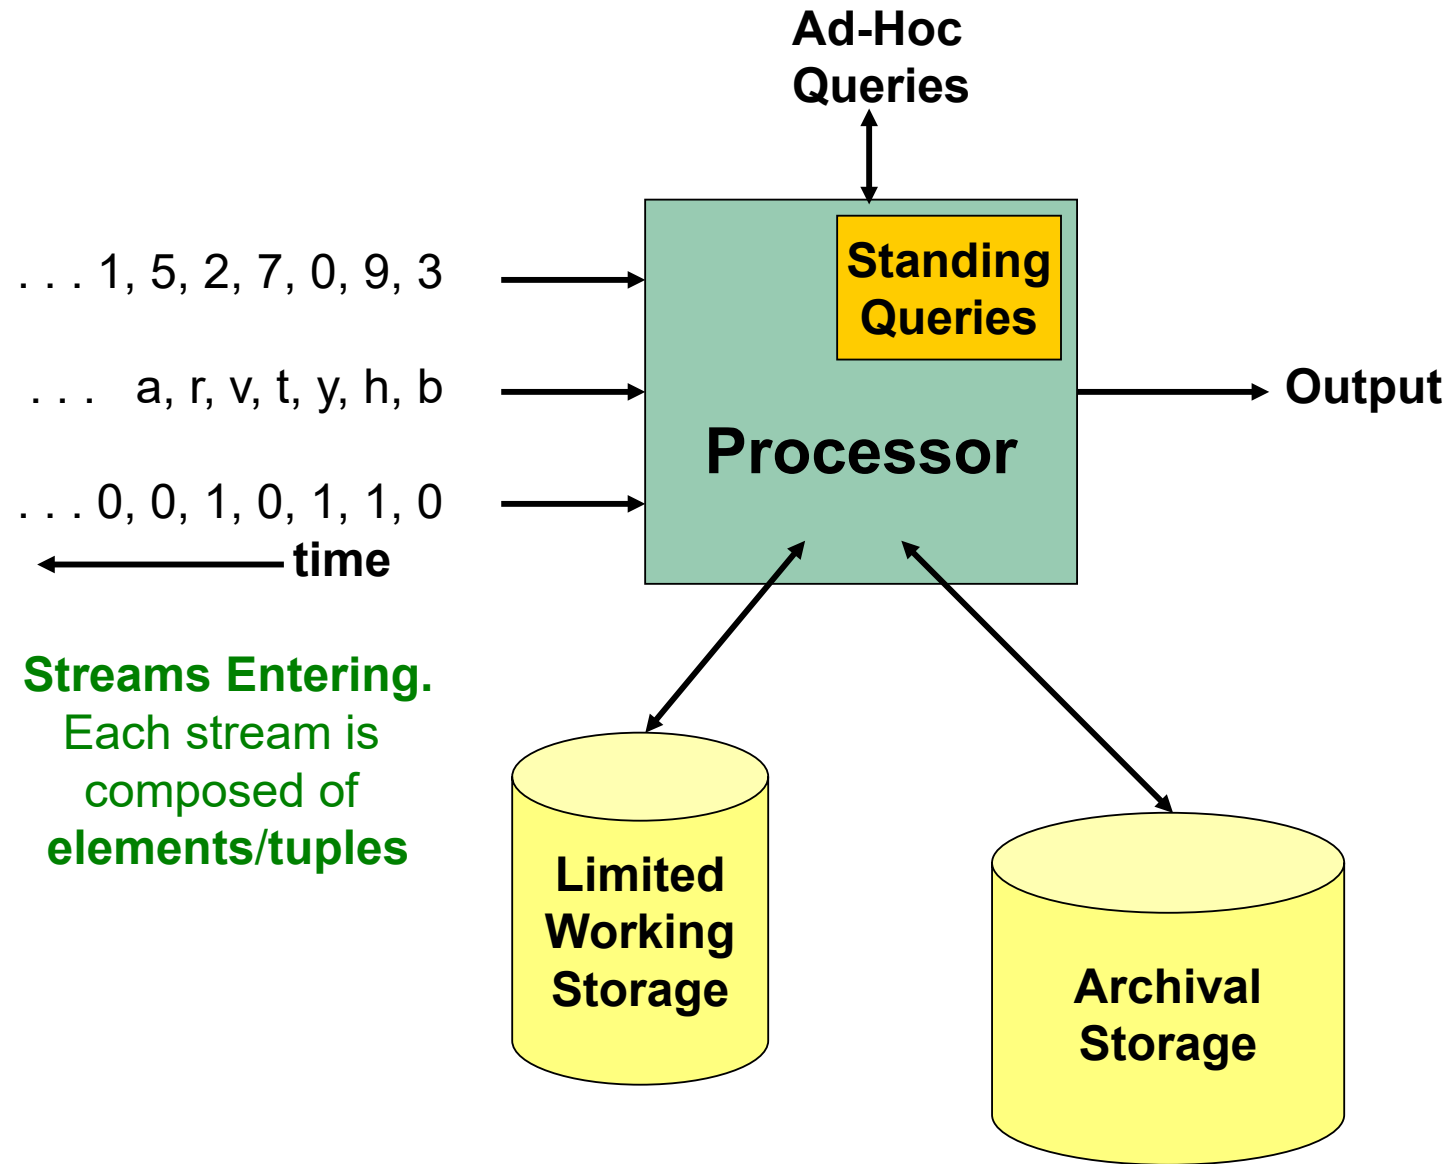
- Input **elements** enter at a **rapid** rate, at one or more input ports
  - We call elements of the stream **tuples**
- The system **cannot** store the **entire** stream
- Q: How do you make critical calculations about the stream using a **limited** amount of memory?

... 1, 5, 2, 7, 0, 9, 3

... a, r, v, t, y, h, b

... 0, 0, 1, 0, 1, 1, 0  
time

# General Stream Processing Model



It is better to use a crude approximation and know the truth, plus or minus 10 percent, than demand an exact solution and know nothing at all.

—Arthur Bloch, The Complete Murphy's Law

# Applications: Networks

- **Mining network streams**

- Finding **abnormal** patterns in sensor reading streams
- Filtering out **spam** calls in phone call streams
- Detect **denial-of-service** attacks in IP packet streams

# Applications: Internet

- **Mining query streams**

- Google wants to know what **queries** are more **frequent** today than yesterday

- **Mining click streams**

- Bytedance wants to know which of its pages are getting an **unusual** number of **hits** in the past hour

- **Mining social network news feeds**

- E.g., look for **trending topics** on Weibo

# Problems on Data Streams

- Types of queries one wants on answer on a data stream (element):
  - **Sampling data from a stream**
    - Construct a random sample
  - **Filtering a data stream**
    - Select elements with property  $x$  from the stream

# Problems on Data Streams

- Types of queries one wants on answer on a data stream (statistics):
  - **Queries over sliding windows**
    - Number of items of **type  $x$**  in the last  $k$  elements of the stream
  - **Counting distinct elements**
    - Number of **distinct** elements in the last  $k$  elements of the stream
  - **Finding frequent elements**
    - Estimate the most **frequent** elements of the last  $k$  elements
  - **Estimating moments**
    - Estimate **avg./std. dev.** of last  $k$  elements



# Sampling from a Data Stream:

## Sampling a fixed-size sample

# Maintaining a fixed-size sample

- Suppose we need to maintain a **random sample**  $S$  of size exactly  $s$  tuples
  - E.g., main memory size constraint
- Suppose at time  $n$  we have seen  $n$  items
  - Each item is in the sample  $S$  with **equal prob.**  $s/n$

**How to think about the problem: say  $s = 2$**

**Stream:** a x c y z k c d e g...

At  $n = 5$ , each of the first 5 tuples is included in the sample  $S$  with equal prob.

At  $n = 7$ , each of the first 7 tuples is included in the sample  $S$  with equal prob.

# Solution: Fixed Size Sample

- **Algorithm**

Store all the first  $s$  elements of the stream to  $\mathcal{S}$

- Suppose we have seen  $n-1$  elements, and now the  $n^{th}$  element arrives ( $n > s$ )
  - With probability  $s/n$ , keep the  $n^{th}$  element, else discard it
  - If we picked the  $n^{th}$  element, then it replaces one of the  $s$  elements in the sample  $\mathcal{S}$ , picked uniformly at random

- This algorithm maintains a sample  $\mathcal{S}$  with the desired property:
  - After  $n$  elements, the sample contains each element seen so far with probability  $s/n$

# Proof: By Induction

- **We prove this by induction:**

- Assume that after  $n$  elements, the sample contains each element seen so far with probability  $s/n$
- We need to show that after seeing element  $n+1$  the sample maintains the property
  - Sample contains each element seen so far with probability  $s/(n+1)$

- **Base case:**

- After we see  $n=s$  elements the sample **S** has the desired property
  - Each out of  $n=s$  elements is in the sample with probability  $s/s = 1$

# Proof: By Induction

- **Inductive hypothesis:** After  $n$  elements, the sample  $\mathcal{S}$  contains each element seen so far with prob.  $s/n$
- Now element  $n+1$  arrives

- **Inductive step:** For elements already in  $\mathcal{S}$ , probability that the algorithm keeps it in  $\mathcal{S}$  is:

$$\underbrace{\left(1 - \frac{s}{n+1}\right)}_{\text{Element } n+1 \text{ discarded}} + \underbrace{\left(\frac{s}{n+1}\right)}_{\text{Element } n+1 \text{ not discarded}} \underbrace{\left(\frac{s-1}{s}\right)}_{\text{Element in the sample not picked}} = \frac{n}{n+1}$$

- So, at time  $n$ , tuples in  $\mathcal{S}$  were there with prob.  $s/n$
- Time  $n \rightarrow n+1$ , tuple stayed in  $\mathcal{S}$  with prob.  $n/(n+1)$
- So prob. tuple is in  $\mathcal{S}$  at time  $n+1 = \frac{s}{n} \cdot \frac{n}{n+1} = \frac{s}{n+1}$

# Filtering Data Streams

# Applications

- Email **spam filtering**
  - We know 1 billion “**good**” email addresses
  - If an email comes from one of these, it is **NOT** spam
- **Publish-subscribe** systems
  - You are collecting lots of messages
  - People express interest in certain sets of **keywords**
  - Determine whether each message matches user’s **interest**

# Filtering Data Streams

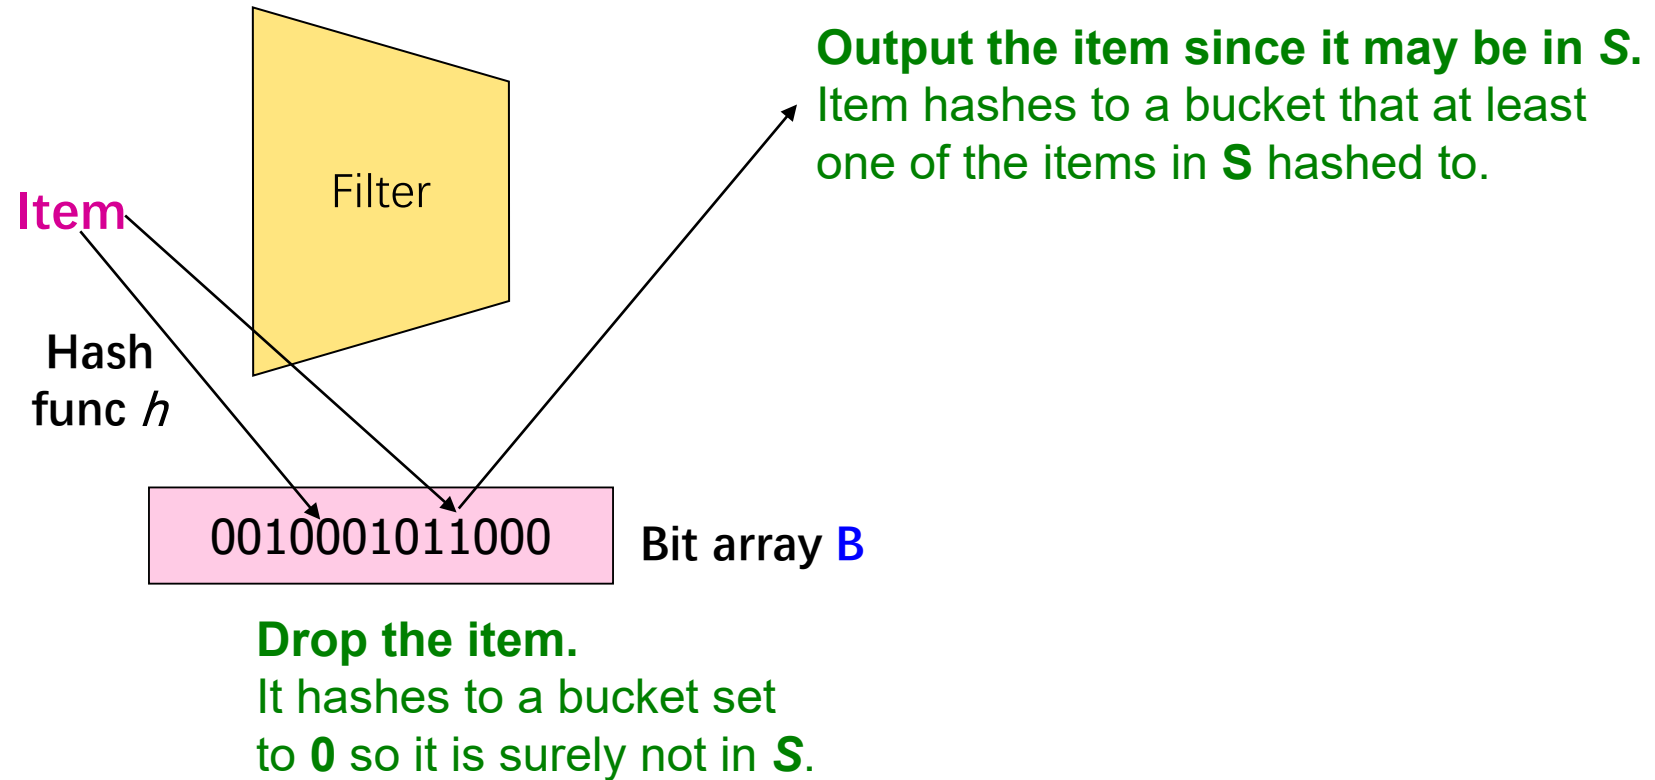
- Each element of data stream is a tuple
- Given a list of keys  $S=[key_1, key_2, \dots]$
- **Determine which tuples of stream are in  $S$**
- Obvious solution: store and compare
  - But suppose we **do not have enough memory** to store all of  $S$
  - The **complexity** is  $O(S)$ , which can be big.



# First Cut Solution

- Given a set of keys  $S$  that we want to filter
- Create a **bit array**  $B$  of  $n$  bits, initially all  $0$ s
- Choose a **hash function**  $h$  with range  $[0, n)$
- Hash each member of  $s \in S$  to one of  $n$  buckets, and set that bit to  $1$ , i.e.,  $B[h(s)] = 1$
- Hash each element  $a$  of the stream and output only those that hash to bit that was set to  $1$ 
  - Output  $a$  if  $B[h(a)] = 1$

# First Cut Solution



- Creates **false positives** but **no false negatives**
  - If the item is in  $S$  we surely output it, if not we may still output it

# First Cut Solution


- $|S| = 1$  billion email addresses  
 $|B| = 1\text{GB} = 8$  billion bits, for the hash values
- If the email address is in  $S$ , then it surely hashes to a bucket that has the bit set to **1**, so it always gets through (*no false negatives*)
- Approximately **1/8** of the bits are set to **1**, so about **1/8** of the addresses not in  $S$  get through to the output (*false positives*)

# Analysis: Throwing Darts

- More accurate analysis for the number of **false positives**
- Consider: If we throw  $m$  darts into  $n$  equally likely targets, what is the probability that a target gets at least one dart?
- **In our case:**
  - **Targets** = bits/buckets
  - **Darts** = hash values of items

# Analysis: Throwing Darts

- We have  $m$  darts,  $n$  targets
- What is the probability that **a target gets at least one dart**?



# Analysis: Throwing Darts

- Fraction of 1s in the array B  
= probability of false positive =  $1 - e^{-m/n}$
- **Example:**  $10^9$  darts,  $8 \cdot 10^9$  targets
  - Fraction of 1s in B =  $1 - e^{-1/8} = 0.1175$ 
    - Compare with our earlier estimate:  $1/8 = 0.125$
- How to further **improve** this false positive probability?
- Similar to LSH: Bloom Filter.

# Bloom Filter

- Consider:  $|\mathbf{S}| = m$ ,  $|\mathbf{B}| = n$
- Use  $k$  independent hash functions  $h_1, \dots, h_k$
- **Initialization:**
  - Set  $\mathbf{B}$  to all **0s**
  - Hash each element  $s \in \mathbf{S}$  using each hash function  $h_i$ , set  $\mathbf{B}[h_i(s)] = \mathbf{1}$  (for each  $i = 1, \dots, k$ )
- **Run-time:**
  - When a stream element with key  $x$  arrives
    - If  $\mathbf{B}[h_i(x)] = \mathbf{1}$  for all  $i = 1, \dots, k$  then declare that  $x$  is in  $\mathbf{S}$ 
      - That is,  $x$  hashes to a bucket set to **1** for every hash function  $h_i(x)$
    - Otherwise discard the element  $x$

# Bloom Filter — Analysis

- **What fraction of the bit vector  $B$  are 1s?**
  - Throwing  $k \cdot m$  darts at  $n$  targets
  - So fraction of 1s is  $(1 - e^{-km/n})$  (false positive of 1 hash function)
- But we have  $k$  independent hash functions and we only let the element  $x$  through **if all**  $k$  hash element  $x$  to a bucket of value **1**
- So, **false positive probability** =  $(1 - e^{-km/n})^k$




# Bloom Filter – Analysis

- $m = 1$  billion,  $n = 8$  billion

- $k = 1$ :  $(1 - e^{-1/8}) = 0.1175$
- $k = 2$ :  $(1 - e^{-1/4})^2 = 0.0493$

- What happens as we keep increasing  $k$ ?

- “Optimal” value of  $k$ :  $n/m \ln(2)$ 
  - In our case: Optimal  $k = 8 \ln(2) = 5.54 \approx 6$ 
    - Error at  $k = 6$ :  $(1 - e^{-1/6})^2 = 0.0235$



# Bloom Filter: Wrap-up

- Bloom filters guarantee **no false negatives**, and use limited memory
  - Great for **pre-processing** before more expensive checks
- Suitable for **hardware** implementation
  - Hash function computations can be **parallelized**
- Is it better to have **1 big B** or **k small Bs**?
  - **It is the same:**  $(1 - e^{-km/n})^k$  vs.  $(1 - e^{-m/(n/k)})^k$
  - But keeping **1 big B** is simpler
- Disadvantage: only insertion, no deletion from Bloom Filter.


# Count-Min Sketch

# Count Element Frequency

- Faced with big data streams, storing all elements and corresponding frequencies is **impossible**.
- **Approximate** counts are acceptable.
- We can use **hashing** again.


# Approximate Counts with Hashing

- **Initialization:**  $count[i] = 0$ , for  $i \in [1, w]$
- **Increment** count of element  $a$ :  $count[h(a)] += 1$
- **Retrieve** count of element  $a$ :  $count[h(a)]$



# Improvement: More Hash Functions

- We use  $d$  pairwise independent hash functions
- **Increment** count of element  $a$ :  $count[i, h_i(a)] += 1$  for  $i \in [1, d]$
- **Retrieve** count of element  $a$ :  $\min_{i \in [1, d]} count[i, h_i(a)]$



# Guarantees

- Theorem[1]: with probability  $1 - \delta$ , the error is at most  $\varepsilon * \text{count}$ . Concrete values for these error bounds can be chosen by setting  $w = \left\lceil \frac{e}{\varepsilon} \right\rceil$  and  $d = \left\lceil \ln\left(\frac{1}{\delta}\right) \right\rceil$ ,  $e \approx 2.718$ .
  - Adding another **hash** function **exponentially** decreases the chance of hash conflicts
  - Increasing the **width** helps spread up the counts with a **linear** effect

# Queries over a Sliding Window



# Sliding Windows

- A useful model of stream processing is that queries are **within** a **window** of length  $N$  – the  $N$  most recent elements received
  - **Amazon example:** For every product  $X$  we keep **0/1 stream** of whether that product was **sold** in the  **$n$ -th transaction**.  
We want answer queries, **how many times we sold  $X$  in the last  $k$  sales**.



Suppose we keep a window with length  $N=6$ ,  
we can query on the last  $k$  transactions, for  $k \leq N$ .

# Counting Bits over a Sliding Window

- **Problem:**


- Given a stream of **0s** and **1s**
- **How many 1s are in the last  $k$  bits?** where  $k \leq N$

- **Obvious solution:** Store the most recent  $N$  bits

- When new bit comes in, discard the  $N+1^{\text{st}}$  bit
- **Not feasible** when  $N$  is so **large** that the data cannot be stored in memory and cannot answer in short time

- **Approximation solution?**

- **Without any assumptions on the data distribution**



# DGIM Method

## DGIM(Datar-Gionis-Indyk-Motwani) Algorithm

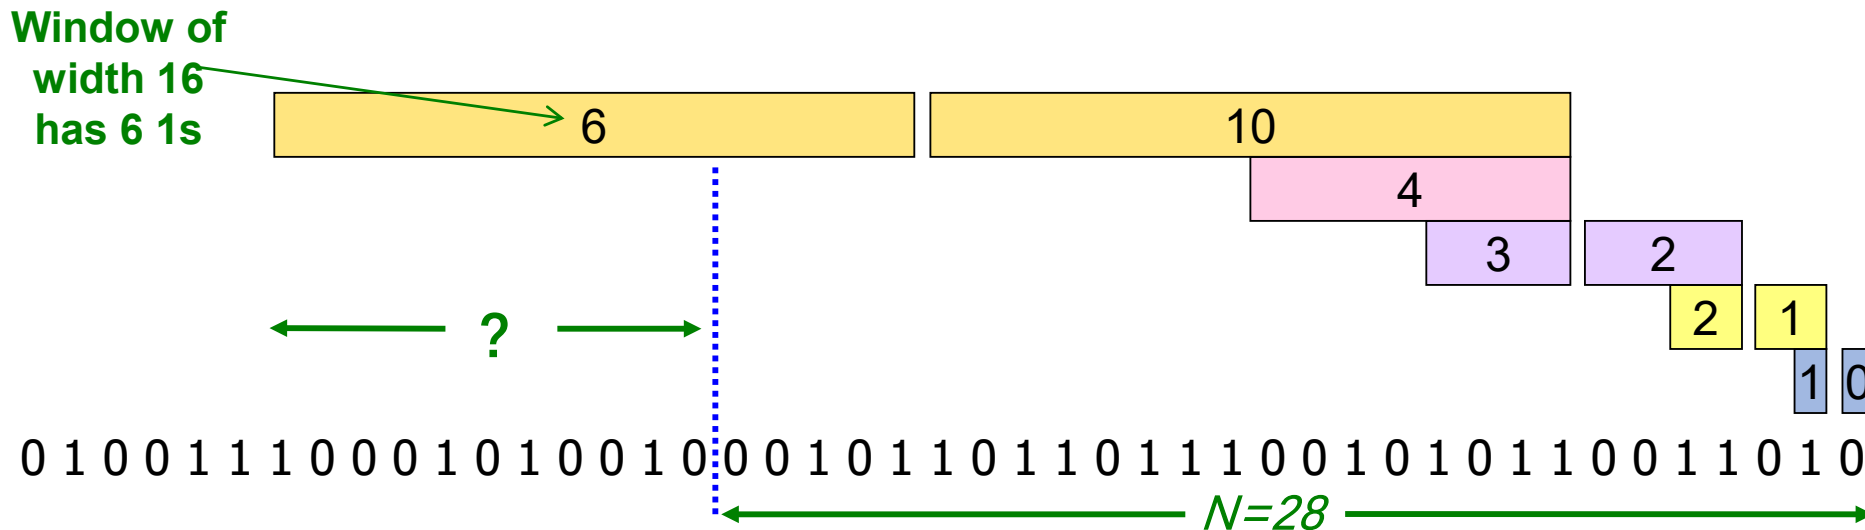
- Does **not** have assumptions on data distribution
- Only stores  $O(\log^2 N)$  bits per stream
- Solution gives **approximate** answer, **never off** by more than **50%**
  - Error factor can be reduced to any fraction  $> 0$ , with more complicated algorithm and proportionally more stored bits

# Idea: Exponential Windows

- **First trial:**

- Summarize **exponentially increasing** regions of the stream, looking backward, to answer queries over last  $k$  items ( $k \leq N$ ).
- Drop small regions if there are more than two on the same level(keep the leftmost)

1. when a bit comes in, create a bucket of length 1 with the proper count (0 or 1).
2. If any level has 3 buckets:
  - a) add the rightmost two and create a bucket at the next higher level (twice the length) with that sum.
  - b) delete the leftmost two buckets, keeping only the rightmost of the three.
3. Repeat (2) recursively for progressively higher levels.



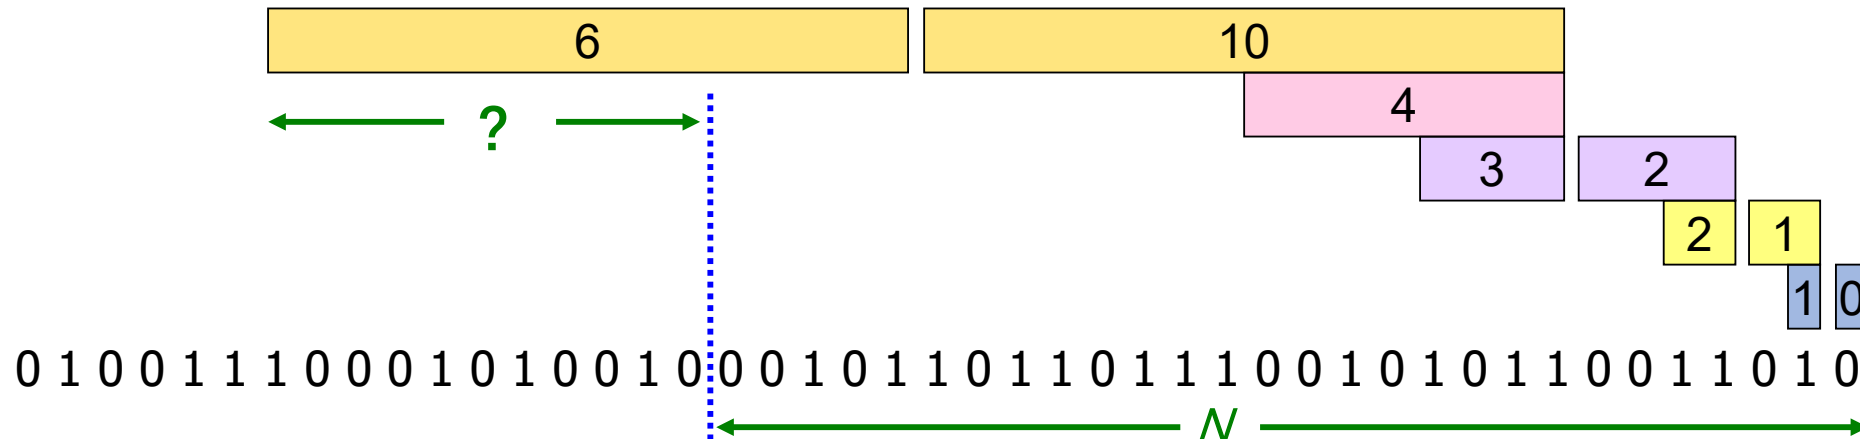
We can reconstruct the count of the last  $N$  bits, except we are not sure how many of the last **6 1s** are included in the  $N$

# What's Good?

- Stores only  $O(\log^2 N)$  bits
  - $O(\log N)$  counts of  $\log_2 N$  bits each
- Easy update as more bits enter
- Error in count no greater than the number of **1s** in the “unknown” area


# What's Not So Good?

- **The relative error could be unbounded!**
  - **Relative error = error / true count**
  - Consider the case that all the **1s** are in the unknown area(?) part) and the rest are all 0s. Here the relative error is infinite.



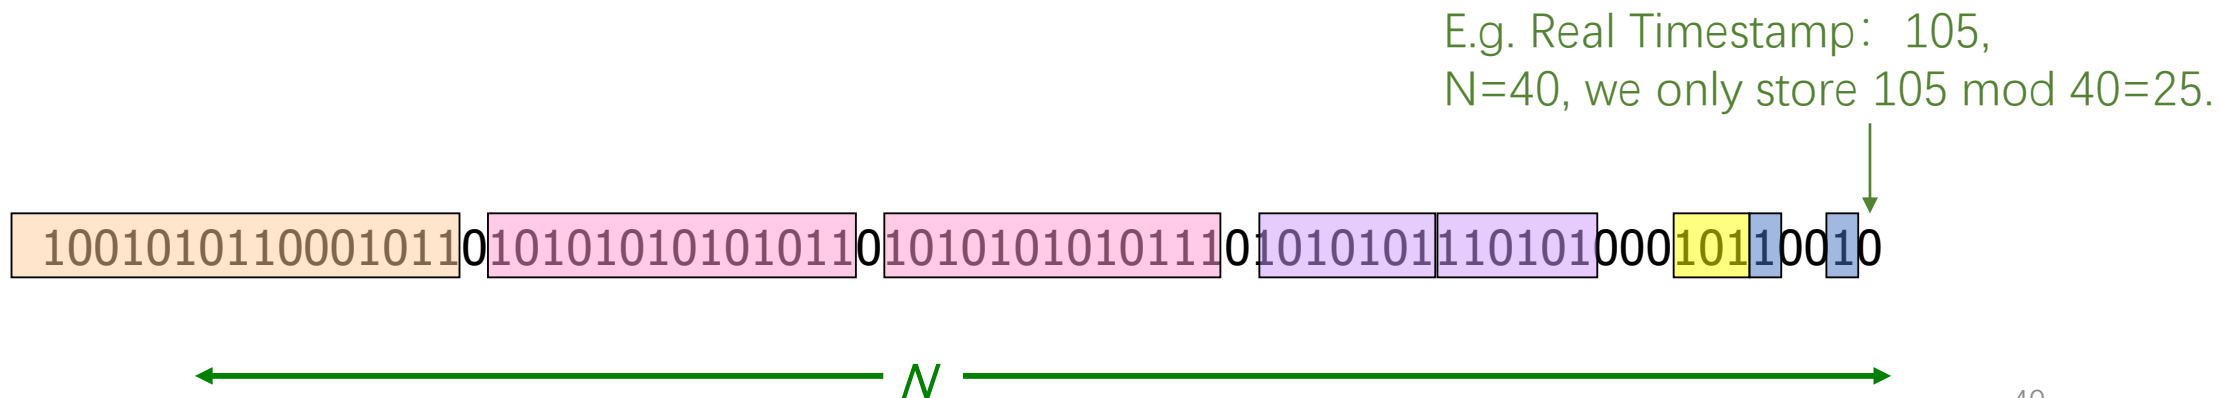
# Fixup: DGIM method

- **Idea:** Instead of summarizing fixed-length blocks, summarize blocks with specific number of **1s**:
  - Let the block **sizes** (number of **1s**) increase **exponentially**
  - **Data dependent**
- When there are few 1s in the window, block sizes stay small, so errors are small



# DGIM: Timestamps

- Each bit in the stream has a *timestamp*, starting **1, 2, ...**
- Record timestamps **modulo  $N$**  (**the window size**), so we can represent any **relevant** timestamp in  $O(\log_2 N)$  bits

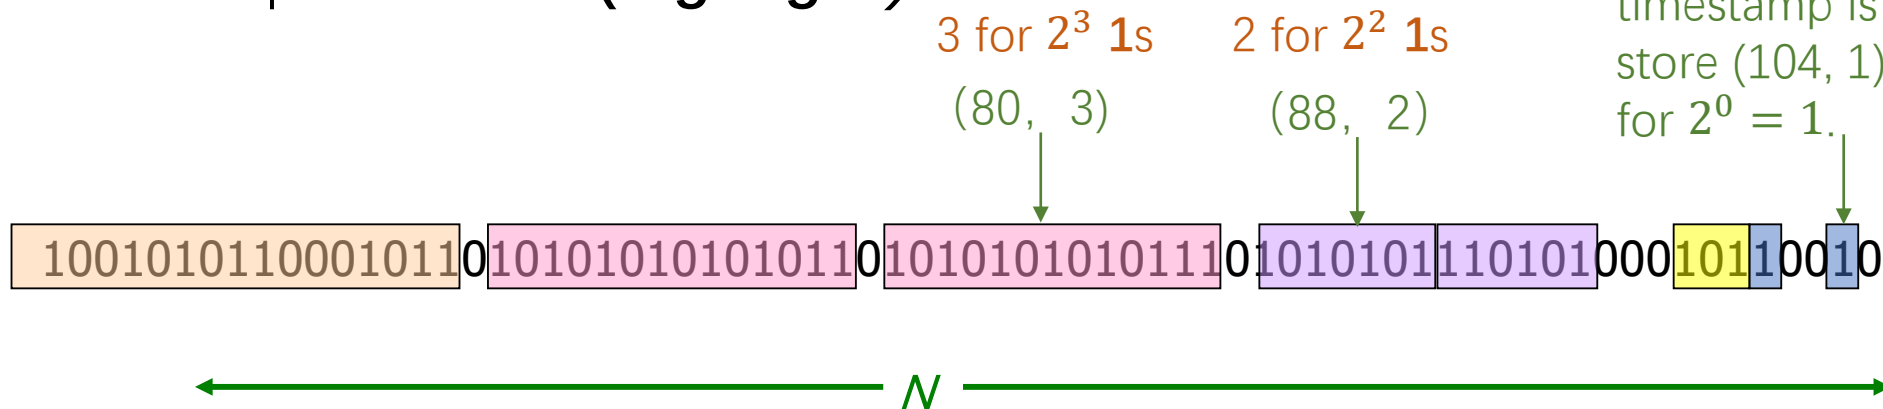




# DGIM: Buckets

- A *bucket* in the DGIM method is a record consisting of:
  - The timestamp of its end [ $O(\log N)$  bits]
  - The number of 1s between its beginning and end [ $O(\log \log N)$  bits]


- **Constraint on buckets:**  
Number of **1s** must be a power of 2
  - That explains the  $O(\log \log N)$



E.g. In this window, if the timestamp of the last timestamp is 105, we actually store (104, 1) here, or (104, 0) for  $2^0 = 1$ .

# Representing a Stream by Buckets

- Either **one** or **two** buckets with the same **power-of-2 number** of **1s**
- Buckets **do not overlap** in timestamps
- Buckets are **sorted** by **size**
  - Earlier buckets are not smaller than later buckets
- Buckets **disappear** when their end-time is **>  $N$**  time units in the past



# Updating Buckets

- When a new bit comes in, **drop** the last (oldest) bucket if its end-time is **prior to  $N$**  time units before the current time
- **2 cases:** Current bit is **0** or **1**
- **If the current bit is 0:**  
**no other changes are needed**

# Updating Buckets

- **If the current bit is 1:**
  - (1) Create a new bucket of size **1**, for just this bit
    - **End timestamp = current time**
  - (2) If there are now three buckets of size 1, combine the oldest two into a bucket of size 2
  - (3) If there are now three buckets of size 2, combine the oldest two into a bucket of size 4
  - (4) Continue until there are at most two buckets of size  $2^i$

# Example: Updating Buckets

**Current state of the stream:**

100101011000101101010101010101101010101010111010101011101010111010100010110010

**Bit of value 1 arrives**

001010110001011010101010101011010101010111010101011101010111010100010110010101

**Two smallest buckets get merged into a size-2 bucket**

0010101100010110101010101010110101010101110101010111010100010110010101

**Next bit 1 arrives, new size-1 bucket is created, then 0 comes, then 1:**

010110001011010101010101010110101010101110101010111010100010110010110101

**Buckets get merged...**


010110001011010101010101010110101010101110101010111010100010110010110101

**State of the buckets after merging**

010110001011010101010101010110101010101110101010111010100010110010110101


# How to Query?

- To estimate the number of 1s in the most recent  $N$  bits:
  1. Sum the sizes of **all** buckets **but the last partially overlapping with the query**
  2. Add **half** the size of the last bucket
- **Remember:** We do **not** know how many **1s** of the last bucket are still within the wanted window




# Error Bound: Proof

- **Why is error 50%? Let's prove it!**
- Suppose the last bucket has size  $2^r$
- Then by assuming  $2^{r-1}$  (i.e., half) of its **1s** are still within the window, we make an error of at most  $2^{r-1}$
- Since there is at least one bucket of each of the sizes less than  $2^r$ , the true sum is at least  
 $1 + 2 + 4 + \dots + 2^{r-1} = 2^r - 1$
- Thus, **relative error at most 50%**



# Further Reducing the Error

- Instead of maintaining **1** or **2** of each size bucket, **we allow either  $s-1$  or  $s$  buckets ( $s > 2$ )**
  - Except for the largest size buckets, where we can have any number between **1** and  **$s$**  of those.



- **Error is at most  $\frac{2^{r-1}}{(s-1)(2^r-1)} = O(1/s)$**
- By picking  **$s$**  appropriately, we can tradeoff between number of bits we store and the error



# Extensions

- Can we handle the case where the stream is not bits, but **integers**, and we want the **sum** of the last  $k$  elements?
- We want the sum of the last  $k$  elements
  - Amazon: Avg. price of last  $k$  sales
- **Solution:**
  - If you know all have at most  $m$  bits
    - Treat  $m$  bits of each integer as a separate stream
    - Use DGIM to count **1s** in each integer
    - The sum is  $= \sum_{i=0}^{m-1} c_i 2^i$

11111111000000001110101010101101010101010111010101011101010001011001**0**  
11111111000000001110101010101101010100010111010101011101010001111001**1**

Two streams represent 1

# Counting Distinct Elements

# Counting Distinct Elements

- **We often ask questions like:**

- How many **distinct people** visit the website?
- How many **distinct products** have we sold in the last week?

- **Problem:**

- Data stream consists of a universe of **elements** chosen from a **set of size  $N$**
- Maintain a **count** of the number of **distinct elements** seen so far

# Using Small Storage

- **Obvious approach:**  
Maintain the set of elements seen so far
- **Real problem:** What if we do not have space to maintain the set of elements seen so far?
- **Same philosophy as previous:**
  - Estimate the count in an **unbiased** way
  - Accept that the count may have a **little error**, but limit the probability that the error is large

# Flajolet-Martin Approach

- Pick a hash function  $h$  that maps each of the  $N$  elements to at least  $\log_2 N$  bits
- For each stream element  $a$ , let  $r(a)$  be the number of trailing 0s in  $h(a)$ 
  - $r(a)$  = position of first 1 counting from the right
    - E.g., say  $h(a) = 12$ , then 12 is 1100 in binary, so  $r(a) = 2$
- Record  $R = \text{the maximum } r(a) \text{ seen}$ 
  - $R = \max_a r(a)$ , over all the items  $a$  seen so far
- Estimated number of distinct elements =  $2^R$

# Why It Works: Intuition

- **Rough and heuristic intuition:**
  - $h(a)$  hashes  $a$  with **equal prob.** to any of  $N$  values
  - Then  $h(a)$  is a sequence of  $\log_2 N$  bits, where  $2^{-r}$  fraction of all  $a$ s have a tail of  $r$  zeros
    - About 50% of  $a$ s hash to **\*\*\*0**
    - About 25% of  $a$ s hash to **\*\*00**
    - So, if we saw the longest tail of  $r=2$  (i.e., item hash ending **\*100**) then we have probably seen **about 4** distinct items so far
  - **So, it takes to hash about  $2^r$  items before we see one with zero-suffix of length  $r$**

# Why It Works: More formally

- Now we show why Flajolet-Martin works
- Formally, we will show that **probability of finding a tail of  $r$  zeros:**
  - Goes to **1** if  $m \gg 2^r$
  - Goes to **0** if  $m \ll 2^r$where  $m$  is the number of distinct elements seen so far in the stream
- Thus,  $2^R$  will almost always be **around  $m$ !**

# Why It Works: More formally

- What is the probability that a given  $h(a)$  ends in at least  $r$  zeros is  $2^{-r}$ 
  - $h(a)$  hashes elements uniformly at random
  - Probability that a random number ends in at least  $r$  zeros is  $2^{-r}$
- Then, the probability of **NOT** seeing a tail of length  $r$  among  $m$  elements:

$$(1 - 2^{-r})^m$$

Prob. all end in fewer than  $r$  zeros.

Prob. that given  $h(a)$  ends in fewer than  $r$  zeros



# Why It Works: More formally

- **Note:**  $(1 - 2^{-r})^m = (1 - 2^{-r})^{2^r (m2^{-r})} \approx e^{-m2^{-r}}$
- **Prob. of NOT finding a tail of length  $r$  is:**
  - If  $m \ll 2^r$ , then prob. tends to **1**
    - $(1 - 2^{-r})^m \approx e^{-m2^{-r}} = 1$  as  $m/2^r \rightarrow 0$
    - So, the probability of finding a tail of length  $r$  tends to **0**
  - If  $m \gg 2^r$ , then prob. tends to **0**
    - $(1 - 2^{-r})^m \approx e^{-m2^{-r}} = 0$  as  $m/2^r \rightarrow \infty$
    - So, the probability of finding a tail of length  $r$  tends to **1**
- **Thus,  $2^R$  will almost always be around  $m!$**

# Issues to fix

- $E[2^R]$  is actually **infinite**
  - Probability halves when  $R \rightarrow R+1$ , but value doubles
  - Limit the bits of hashing values to  $L$
- The estimation is **biased**
  - Estimated with  $2^R/\Phi$ , where  $\Phi = 0.77351$  is a correction factor.
- Problems of **high variance**. Improve accuracy.
  - Use many hash functions with samples of  $R$
  - Partition your samples into small groups
  - Take the median of groups
  - Then take the average of the medians

# Computing Moments

# Generalization: Moments

- Suppose a stream has elements chosen from a set  $A$  of  $N$  values (say 1 to  $N$ )
- Let  $m_i$  be the number of times value  $i$  occurs in the stream
- The  $k^{\text{th}}$  moment is

$$\sum_{i \in A} (m_i)^k$$

# Special Cases

$$\sum_{i \in A} (m_i)^k$$

- **0<sup>th</sup> moment** = number of **distinct** elements (Flajolet-Martin Approach)
- **1<sup>st</sup> moment** = count of the **numbers** of elements = length of the stream
- **2<sup>nd</sup> moment** = **surprise number S** = a measure of how uneven the distribution is

E.g. **Stream of length 100, 11 distinct values**

- Item counts: **10, 9, 9, 9, 9, 9, 9, 9, 9, 9, 9** **Surprise S = 910**
- Item counts: **90, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1** **Surprise S = 8,110**

# AMS(Alon–Matias–Szegedy) Method

- AMS method works for **all moments**
- Gives an **unbiased estimate**
- We will just concentrate on the **2<sup>nd</sup> moment**  $S$
- We pick and keep track of some variables  $X$ :
  - For each variable  $X$  we store  $X.e/$  and  $X.val$ 
    - $X.e/$  corresponds to the item  $i$
    - $X.val$  corresponds to the **count** of item  $i$
  - Note this requires a count in main memory, so number of  $X$ s is limited
- **Our goal is to compute**  $S = \sum_i m_i^2$

# One Random Variable (X)

- **How to set  $X.val$  and  $X.el$ ?**


- Assume stream has length  $n$
- Pick some **random** time  $t$  ( $t < n$ ) to start, so that any time is **equally likely**
- Let at time  $t$  the stream have item  $i$ . *We set  $X.el = i$*
- Then we maintain count  $c$  ( $X.val = c$ ) of the number of  $i$ s in the stream starting from the chosen time  $t$

- **Then the estimate of the 2<sup>nd</sup> moment ( $\sum_i m_i^2$ ) is:**

$$S = f(X) = n(2 \cdot c - 1)$$

- Note, we will keep track of multiple Xs, ( $X_1, X_2, \dots, X_k$ ) and our final estimate will be  $S = 1/k \sum_j f(X_j)$

# Expectation Analysis



- 2<sup>nd</sup> moment is  $S = \sum_i m_i^2$
- $c_t$  ... number of times item at time  $t$  appears from time  $t$  onwards ( $c_1=m_a$ ,  $c_2=m_a-1$ ,  $c_3=m_b$ )

$$E[f(X)] = \frac{1}{n} \sum_{t=1}^n n(2c_t - 1)$$

$$= \frac{1}{n} \sum_i n (1 + 3 + 5 + \dots + 2m_i - 1)$$

$m_i$  ... total count of item  $i$  in the stream (we are assuming stream has length  $n$ )

Group times by the value seen


Time  $t$  when the last  $i$  is seen ( $c_t=1$ )

Time  $t$  when the penultimate  $i$  is seen ( $c_t=2$ )

Time  $t$  when the first  $i$  is seen ( $c_t=m_i$ )



# Expectation Analysis



- $E[f(X)] = \frac{1}{n} \sum_i n (1 + 3 + 5 + \dots + 2m_i - 1)$ 
  - calculation:  $(1 + 3 + 5 + \dots + 2m_i - 1) = \sum_{i=1}^{m_i} (2i - 1) = 2 \frac{m_i(m_i+1)}{2} - m_i = (m_i)^2$
- Then  $E[f(X)] = \frac{1}{n} \sum_i n (m_i)^2$
- So,  $E[f(X)] = \sum_i (m_i)^2 = S$
- We have the second moment (in expectation)!

# Higher-Order Moments

- For estimating  $k^{\text{th}}$  moment we essentially use the same algorithm but change the estimate:
  - For  $k=2$  we used  $n(2c - 1)$
  - For  $k=3$  we use:  $n(3c^2 - 3c + 1)$  (where  $c=X.\text{val}$ )
- Why?
  - **For  $k=2$ :** Remember we had  $(1 + 3 + 5 + \dots + 2m_i - 1)$  and we showed terms  $2c-1$  (for  $c=1, \dots, m$ ) sum to  $m^2$ 
    - $\sum_{c=1}^m 2c - 1 = \sum_{c=1}^m c^2 - \sum_{c=1}^m (c - 1)^2 = m^2$
    - So:  $2c - 1 = c^2 - (c - 1)^2$
  - **For  $k=3$ :**  $c^3 - (c-1)^3 = 3c^2 - 3c + 1$
- **Generally:** Estimate =  $n(c^k - (c - 1)^k)$

# Combining Samples

- **In practice:**

- Compute  $f(X) = n(2c - 1)$  for as many variables  $X$  as you can fit in memory
- Average them in groups
- Take median of averages

- **Problem: Streams never end**

- We assumed there was a number  $n$ , the number of positions in the stream
- But real streams go on forever, so  $n$  is a variable – the number of inputs seen so far

# Streams Never End: Fixups

(1) The variables  $X$  have  $n$  as a factor – keep  $n$  separately; just hold the count in  $X$

(2) Suppose we can only store  $k$  counts.

We must throw some  $X$ s out as time goes on:

- **Objective:** Each starting time  $t$  is selected with probability  $k/n$
- **Solution: (fixed-size (Reservoir) sampling!)**
  - Choose the first  $k$  times for  $k$  variables
  - When the  $n^{\text{th}}$  element arrives ( $n > k$ ), choose it with probability  $k/n$
  - If you choose it, throw one of the previously stored variables  $X$  out, with equal probability

# Summary of Streaming Algorithms

- Queries
  - Filtering a data stream
  - Queries over a sliding window
  - Counting distinct elements
  - Estimating moments
- Key techniques
  - Hashing functions
  - Approximation with sketch/summarization
  - Theoretical analysis