

EE359 Big Data Mining

Streaming Algorithms

Jiaxin Ding

John Hopcroft Center



上海交通大学
约翰·霍普克罗夫特
计算机科学中心

John Hopcroft Center for Computer Science



Data Streams

- A data stream is a **sequence** of signals used to transmit or receive information that is in the process of being **transmitted**. In many situations, we do not know the entire data set in advance.
 - **Infinite**
 - **Non-stationary**

. . . 1, 5, 2, 7, 0, 9, 3

. . . a, r, v, t, y, h, b

. . . 0, 0, 1, 0, 1, 1, 0
time

The Stream Model

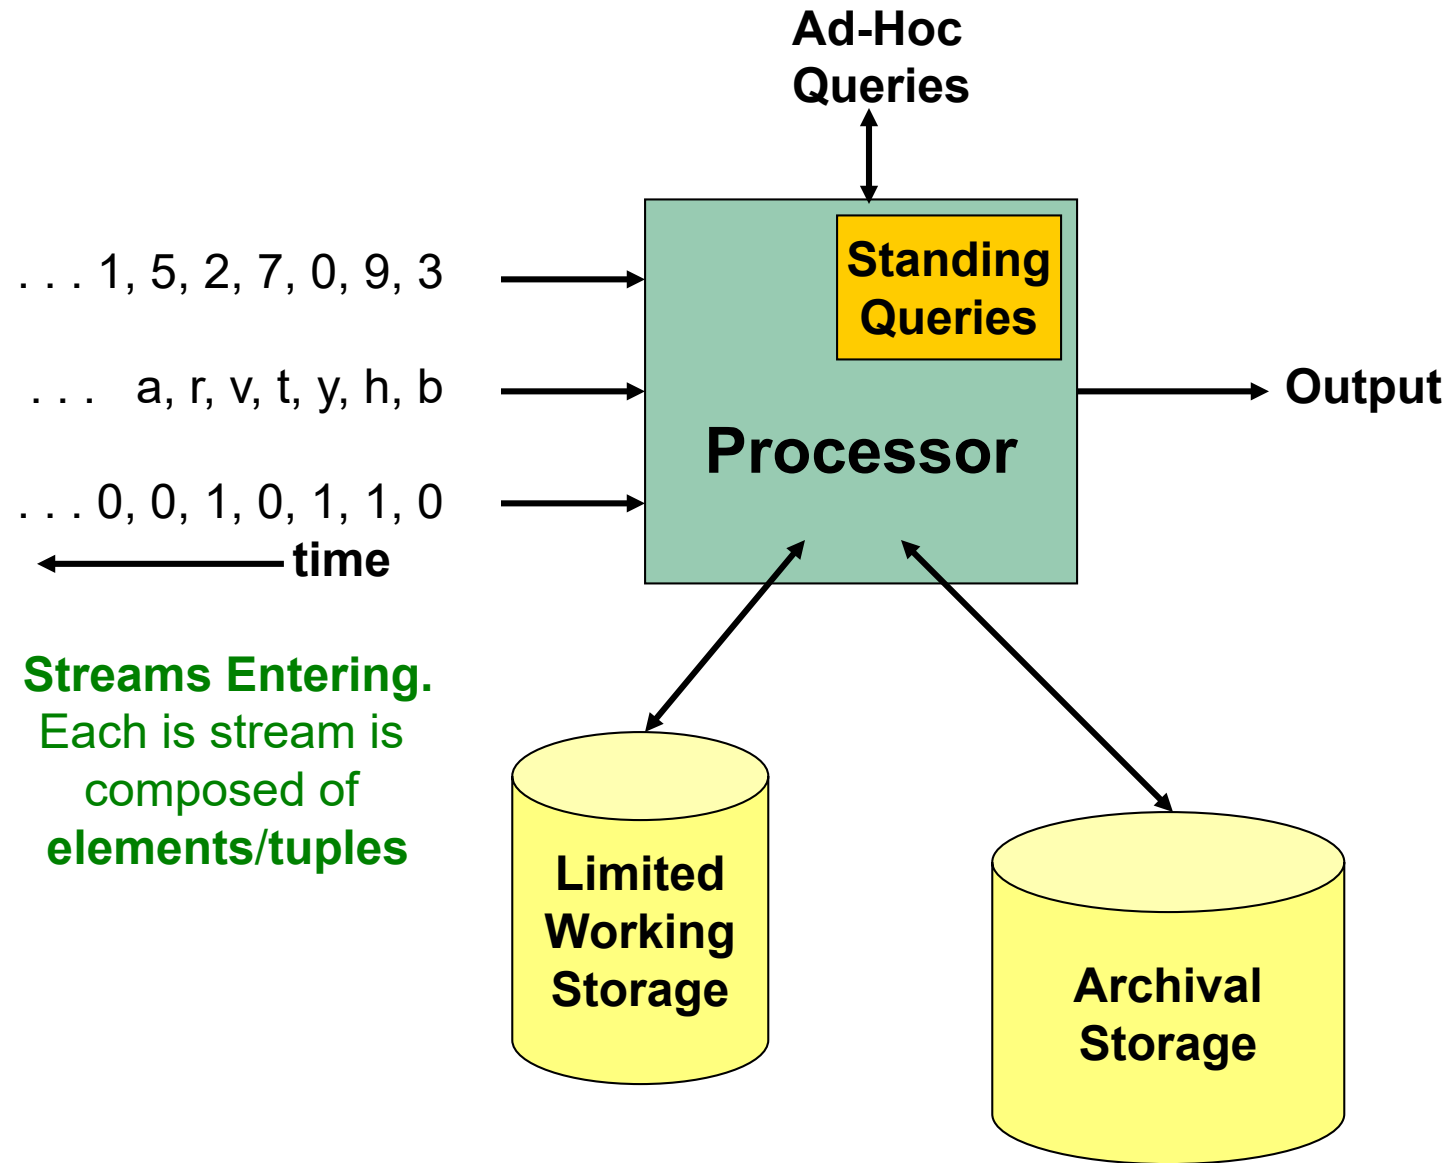
- Input **elements** enter at a **rapid** rate, at one or more input ports
 - We call elements of the stream **tuples**
- The system **cannot** store the **entire** stream
- Q: How do you make critical calculations about the stream using a **limited** amount of memory?

... 1, 5, 2, 7, 0, 9, 3

... a, r, v, t, y, h, b

... 0, 0, 1, 0, 1, 1, 0
time

General Stream Processing Model



Streams Entering.
Each stream is composed of
elements/tuples

It is better to use a crude approximation and know the truth, plus or minus 10 percent, than demand an exact solution and know nothing at all.

—Arthur Bloch, The Complete Murphy's Law

Applications: Networks

- **Mining network streams**

- Finding **abnormal** patterns in sensor reading streams
- Filtering out **spam** calls in phone call streams
- Detect **denial-of-service** attacks in IP packet streams

Applications: Internet

- **Mining query streams**

- Google wants to know what **queries** are more **frequent** today than yesterday

- **Mining click streams**

- Bytedance wants to know which of its pages are getting an **unusual** number of **hits** in the past hour

- **Mining social network news feeds**

- E.g., look for **trending topics** on Weibo

Problems on Data Streams

- Types of queries one wants on answer on a data stream (element):
 - **Sampling data from a stream**
 - Construct a random sample
 - **Filtering a data stream**
 - Select elements with property x from the stream

Problems on Data Streams

- Types of queries one wants on answer on a data stream (statistics):
 - **Queries over sliding windows**
 - Number of items of **type x** in the last k elements of the stream
 - **Counting distinct elements**
 - Number of **distinct** elements in the last k elements of the stream
 - **Finding frequent elements**
 - Estimate the most **frequent** elements of the last k elements
 - **Estimating moments**
 - Estimate **avg./std. dev.** of last k elements

Sampling from a Data Stream:

Sampling a fixed-size sample

Maintaining a fixed-size sample

- Suppose we need to maintain a **random sample** S of size exactly s tuples
 - E.g., main memory size constraint
- Suppose at time n we have seen n items
 - Each item is in the sample S with **equal prob.** s/n

How to think about the problem: say $s = 2$

Stream: a x c y z k c d e g...

At $n = 5$, each of the first 5 tuples is included in the sample S with equal prob.

At $n = 7$, each of the first 7 tuples is included in the sample S with equal prob.

Solution: Fixed Size Sample

- **Algorithm**

Store all the first s elements of the stream to \mathcal{S}

- Suppose we have seen $n-1$ elements, and now the n^{th} element arrives ($n > s$)
 - With probability s/n , keep the n^{th} element, else discard it
 - If we picked the n^{th} element, then it replaces one of the s elements in the sample \mathcal{S} , picked uniformly at random

- This algorithm maintains a sample \mathcal{S} with the desired property:
 - After n elements, the sample contains each element seen so far with probability s/n

Proof: By Induction

- **We prove this by induction:**

- Assume that after n elements, the sample contains each element seen so far with probability s/n
- We need to show that after seeing element $n+1$ the sample maintains the property
 - Sample contains each element seen so far with probability $s/(n+1)$

- **Base case:**

- After we see $n=s$ elements the sample **S** has the desired property
 - Each out of $n=s$ elements is in the sample with probability $s/s = 1$

Proof: By Induction

- **Inductive hypothesis:** After n elements, the sample \mathcal{S} contains each element seen so far with prob. s/n

- Now element $n+1$ arrives

- **Inductive step:** For elements already in \mathcal{S} , probability that the algorithm keeps it in \mathcal{S} is:

$$\underbrace{\left(1 - \frac{s}{n+1}\right)}_{\text{Element } n+1 \text{ discarded}} + \underbrace{\left(\frac{s}{n+1}\right)}_{\text{Element } n+1 \text{ not discarded}} \underbrace{\left(\frac{s-1}{s}\right)}_{\text{Element in the sample not picked}} = \frac{n}{n+1}$$

- So, at time n , tuples in \mathcal{S} were there with prob. s/n
- Time $n \rightarrow n+1$, tuple stayed in \mathcal{S} with prob. $n/(n+1)$
- So prob. tuple is in \mathcal{S} at time $n+1 = \frac{s}{n} \cdot \frac{n}{n+1} = \frac{s}{n+1}$

Filtering Data Streams

Applications

- Email **spam filtering**
 - We know 1 billion “good” email addresses
 - If an email comes from one of these, it is **NOT** spam
- **Publish-subscribe** systems
 - You are collecting lots of messages
 - People express interest in certain sets of **keywords**
 - Determine whether each message matches user’s **interest**

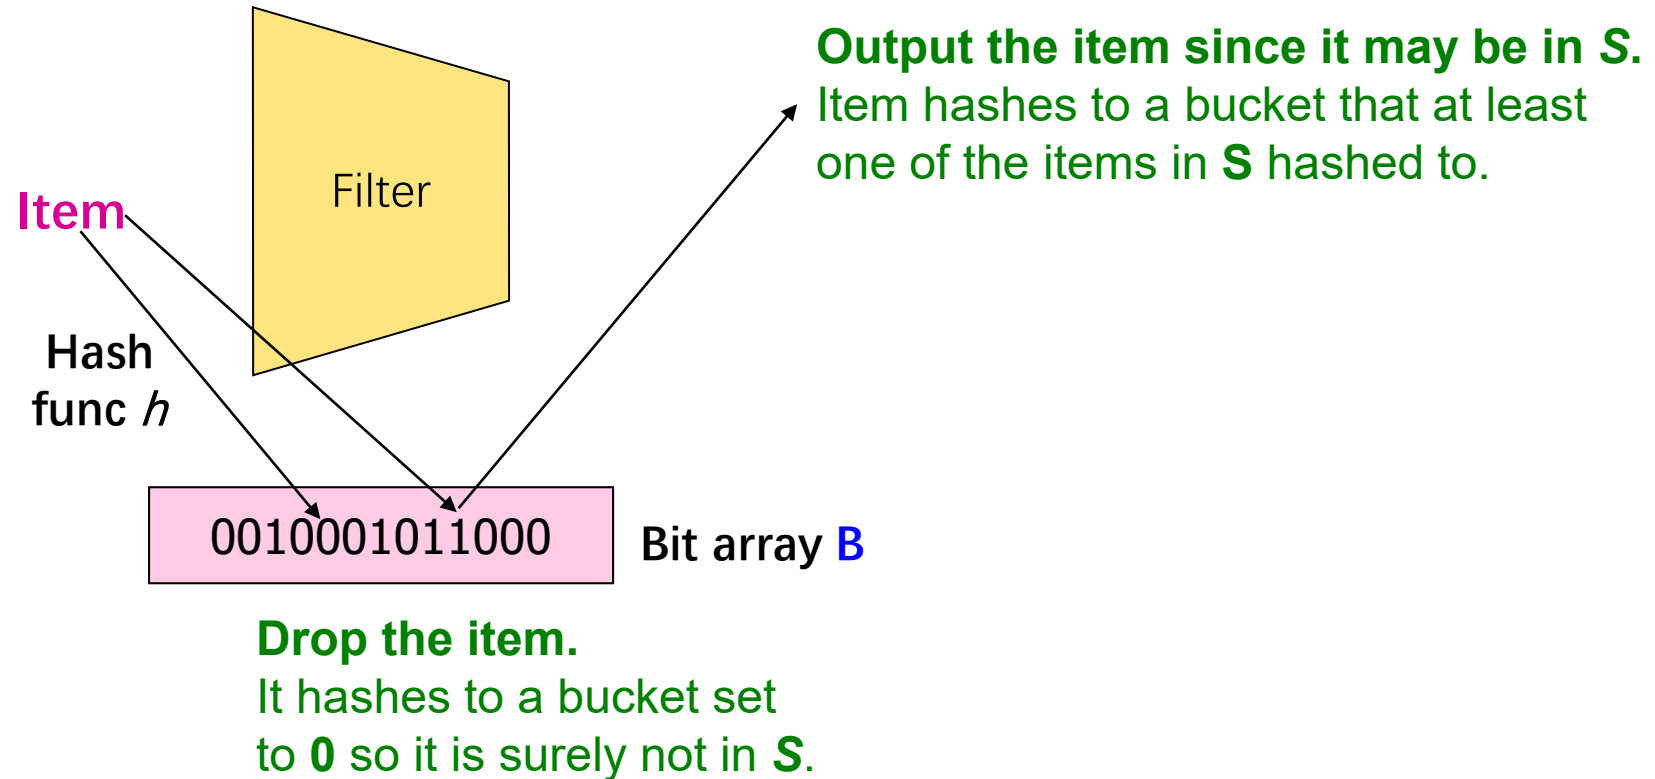
Filtering Data Streams

- Each element of data stream is a tuple
- Given a list of keys $S=[key_1, key_2, \dots]$
- **Determine which tuples of stream are in S**
- Obvious solution: store and compare
 - But suppose we **do not have enough memory** to store all of S
 - The **complexity** is $O(S)$, which can be big.

First Cut Solution

- Given a set of keys S that we want to filter
- Create a **bit array** B of n bits, initially all 0 s
- Choose a **hash function** h with range $[0, n)$
- Hash each member of $s \in S$ to one of n buckets, and set that bit to 1 , i.e., $B[h(s)] = 1$
- Hash each element a of the stream and output only those that hash to bit that was set to 1
 - Output a if $B[h(a)] = 1$

First Cut Solution



- Creates **false positives** but **no false negatives**
 - If the item is in **S** we surely output it, if not we may still output it

First Cut Solution


- $|S| = 1$ billion email addresses
 $|B| = 1\text{GB} = 8$ billion bits, for the hash values
- If the email address is in S , then it surely hashes to a bucket that has the bit set to **1**, so it always gets through (*no false negatives*)
- Approximately **1/8** of the bits are set to **1**, so about **1/8** of the addresses not in S get through to the output (*false positives*)

Analysis: Throwing Darts

- More accurate analysis for the number of **false positives**
- Consider: If we throw m darts into n equally likely targets, what is the probability that a target gets at least one dart?
- **In our case:**
 - **Targets** = bits/buckets
 - **Darts** = hash values of items

Analysis: Throwing Darts

- We have m darts, n targets
- What is the probability that **a target gets at least one dart**?



Analysis: Throwing Darts

- Fraction of 1s in the array B
= probability of false positive = $1 - e^{-m/n}$
- **Example:** 10^9 darts, $8 \cdot 10^9$ targets
 - Fraction of 1s in B = $1 - e^{-1/8} = 0.1175$
 - Compare with our earlier estimate: $1/8 = 0.125$
- How to further **improve** this false positive probability?
- Similar to LSH: Bloom Filter.

Bloom Filter

- Consider: $|\mathbf{S}| = m$, $|\mathbf{B}| = n$
- Use k independent hash functions h_1, \dots, h_k
- **Initialization:**
 - Set \mathbf{B} to all 0 s
 - Hash each element $s \in \mathbf{S}$ using each hash function h_i , set $\mathbf{B}[h_i(s)] = 1$ (for each $i = 1, \dots, k$)
- **Run-time:**
 - When a stream element with key x arrives
 - If $\mathbf{B}[h_i(x)] = 1$ for all $i = 1, \dots, k$ then declare that x is in \mathbf{S}
 - That is, x hashes to a bucket set to 1 for every hash function $h_i(x)$
 - Otherwise discard the element x

Bloom Filter — Analysis

- **What fraction of the bit vector B are 1s?**
 - Throwing $k \cdot m$ darts at n targets
 - So fraction of 1s is $(1 - e^{-km/n})$ (false positive of 1 hash function)
- But we have k independent hash functions and we only let the element x through **if all** k hash element x to a bucket of value **1**
- So, **false positive probability** = $(1 - e^{-km/n})^k$


Bloom Filter – Analysis

- $m = 1$ billion, $n = 8$ billion

- $k = 1$: $(1 - e^{-1/8}) = 0.1175$
- $k = 2$: $(1 - e^{-1/4})^2 = 0.0493$

- What happens as we keep increasing k ?

- “Optimal” value of k : $n/m \ln(2)$
 - In our case: Optimal $k = 8 \ln(2) = 5.54 \approx 6$
 - Error at $k = 6$: $(1 - e^{-1/6})^2 = 0.0235$



Bloom Filter: Wrap-up

- Bloom filters guarantee **no false negatives**, and use limited memory
 - Great for **pre-processing** before more expensive checks
- Suitable for **hardware** implementation
 - Hash function computations can be **parallelized**
- Is it better to have **1 big B** or **k small Bs**?
 - **It is the same:** $(1 - e^{-km/n})^k$ vs. $(1 - e^{-m/(n/k)})^k$
 - But keeping **1 big B** is simpler
- Disadvantage: only insertion, no deletion from Bloom Filter.


Count-Min Sketch

Count Element Frequency

- Faced with big data streams, storing all elements and corresponding frequencies is **impossible**.
- **Approximate** counts are acceptable.
- We can use **hashing** again.

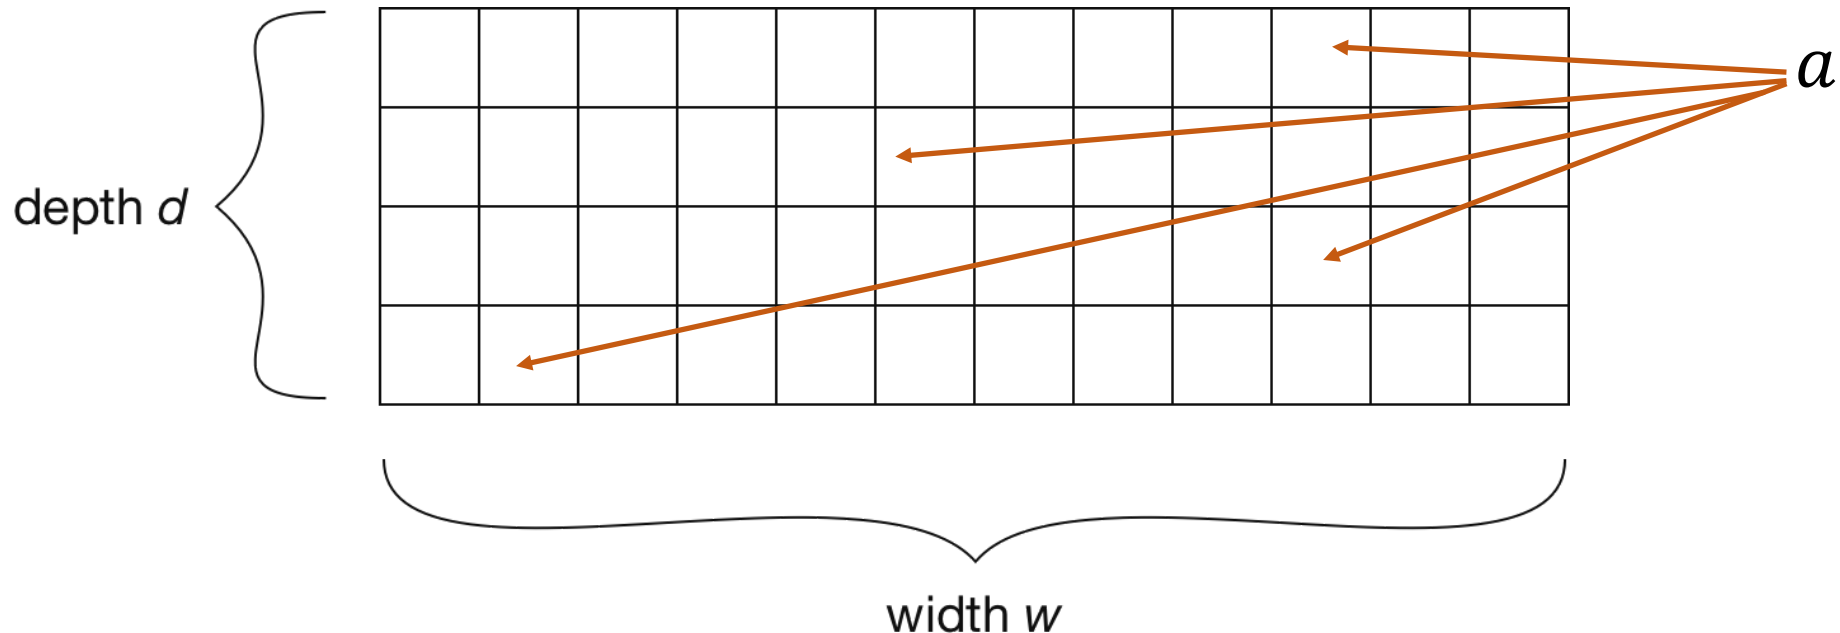
Approximate Counts with Hashing

- **Initialization**: $count[i] = 0$, for $i \in [1, w]$
- **Increment** count of element a : $count[h(a)] += 1$
- **Retrieve** count of element a : $count[h(a)]$



Improvement: More Hash Functions

- We use d pairwise independent hash functions
- **Increment** count of element a : $count[i, h_i(a)] += 1$ for $i \in [1, d]$
- **Retrieve** count of element a : $\min_{i \in [1, d]} count[i, h_i(a)]$




Guarantees

- Theorem[1]: with probability $1 - \delta$, the error is at most $\varepsilon * \text{count}$. Concrete values for these error bounds can be chosen by setting $w = \left\lceil \frac{e}{\varepsilon} \right\rceil$ and $d = \left\lceil \ln\left(\frac{1}{\delta}\right) \right\rceil$, $e \approx 2.718$.
 - Adding another **hash** function **exponentially** decreases the chance of hash conflicts
 - Increasing the **width** helps spread up the counts with a **linear** effect

Queries over a Sliding Window

Sliding Windows

- A useful model of stream processing is that queries are **within** a **window** of length N – the N most recent elements received
 - **Amazon example:** For every product X we keep 0/1 stream of whether that product was **sold** in the n -th **transaction**. We want answer queries, **how many times we sold X in the last k sales**.



Suppose we keep a window with length $N=6$, we can query on the last k transactions, for $k \leq N$.

Counting Bits


- **Problem:**

- Given a stream of **0**s and **1**s
- Be prepared to answer queries of the form
How many 1s are in the last k bits? where $k \leq N$

- **Obvious solution:**

Store the most recent N bits

- When new bit comes in, discard the $N+1^{\text{st}}$ bit
- **Not feasible** when N is so **large** that the data cannot be stored in memory, or even on disk




Counting Bits

- **Real Problem:**

What if we cannot afford to store or compute N bits?


- **E.g.**, we're processing 1 billion streams and $N = 1$ billion



- But we are happy with an **approximate** answer

An attempt: Simple solution

- **Q: How many 1s are in the last N bits?**
- A simple solution that does **not really** solve our problem:
Uniformity assumption



- **Maintain 2 counters:**
 - S : number of **1**s from the beginning of the stream
 - Z : number of **0**s from the beginning of the stream
- How many 1s are in the last N bits? $N \cdot \frac{S}{S+Z}$
- But, what if stream is **non-uniform**?
 - What if distribution changes over time? This is always true in reality.

DGIM Method

- **DGIM**(*Datar-Gionis-Indyk-Motwani Algorithm*) solution that **does not** assume uniformity
- We store $O(\log^2 N)$ bits per stream
- Solution gives **approximate** answer, **never off** by more than **50%**
 - Error factor can be reduced to any fraction > 0 , with more complicated algorithm and proportionally more stored bits

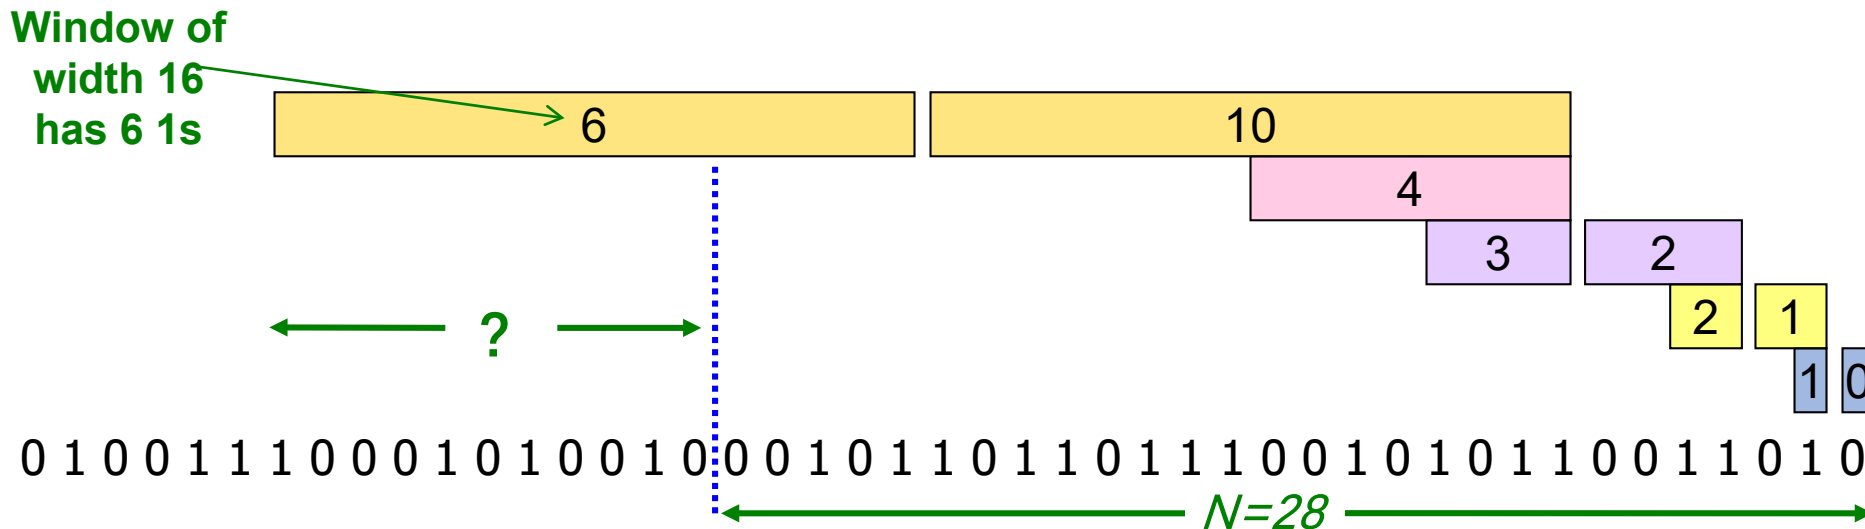
Guess: how to achieve $O(\log^2 N)$ bits to answer queries over the last k items.

Idea: Exponential Windows

- **First trial:**

- Summarize **exponentially increasing** regions of the stream, looking backward, to answer queries over last k items ($k \leq N$).
- Drop small regions if there are more than two on the same level (keep the leftmost)

1. when a bit comes in, create a bucket of length 1 with the proper count (0 or 1).
2. If any level has 3 buckets:
 - a) add the rightmost two and create a bucket at the next higher level (twice the length) with that sum.
 - b) delete the leftmost two buckets, keeping only the rightmost of the three.
3. Repeat (2) recursively for progressively higher levels.



We can reconstruct the count of the last N bits, except we are not sure how many of the last **6 1s** are included in the N

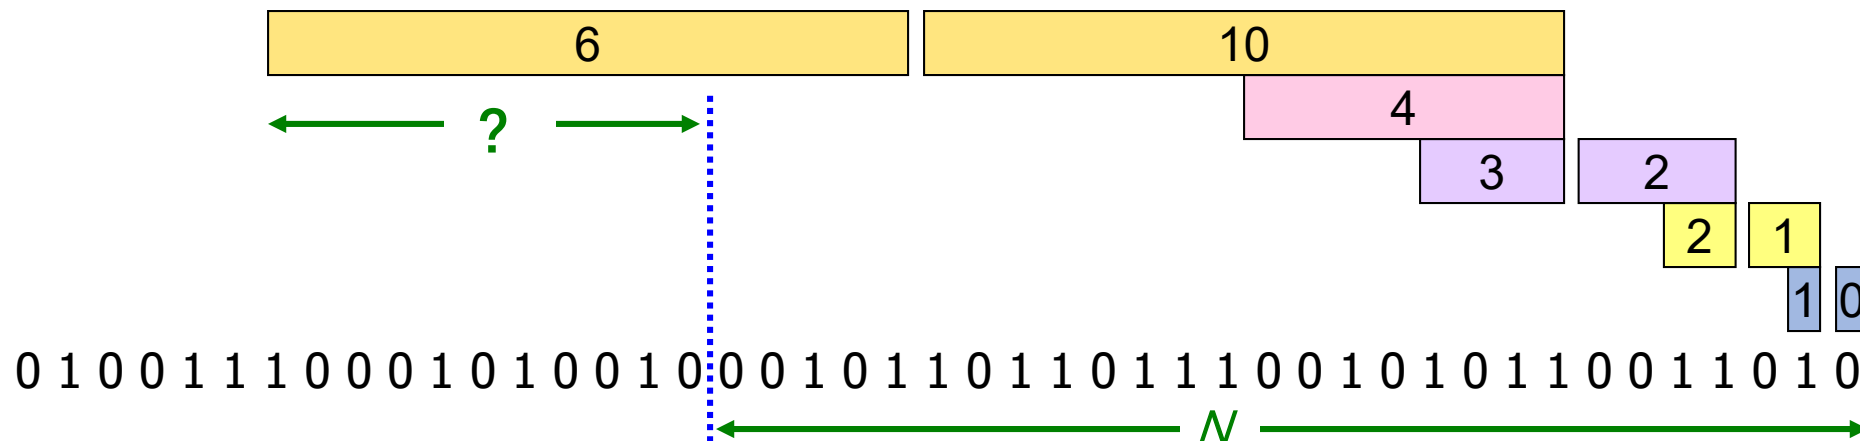
What's Good?

- Stores only $O(\log^2 N)$ bits
 - $O(\log N)$ counts of $\log_2 N$ bits each
- Easy update as more bits enter
- Error in count no greater than the number of **1s** in the “unknown” area

What's Not So Good?


Relative error = error / true count

- As long as the **1s** are fairly **evenly distributed**, the error due to the unknown region is small
- **But the relative error could be unbounded!**
 - Consider the case that all the **1s** are in the unknown area(?) part) and the rest are all 0s. Here the relative error is infinite.



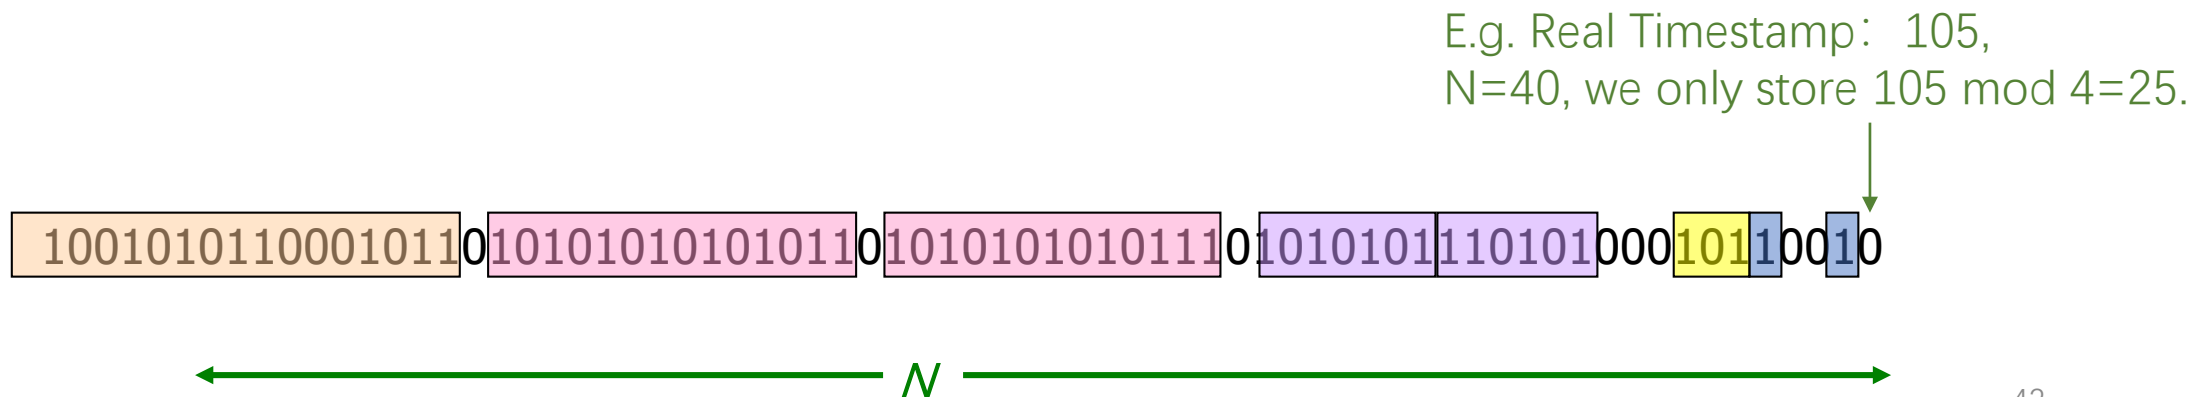
Fixup: DGIM method

- **Idea:** Instead of summarizing fixed-length blocks, summarize blocks with specific number of **1s**:
 - Let the block **sizes** (number of **1s**) increase **exponentially**
 - **Data dependent**
- When there are few 1s in the window, block sizes stay small, so errors are small



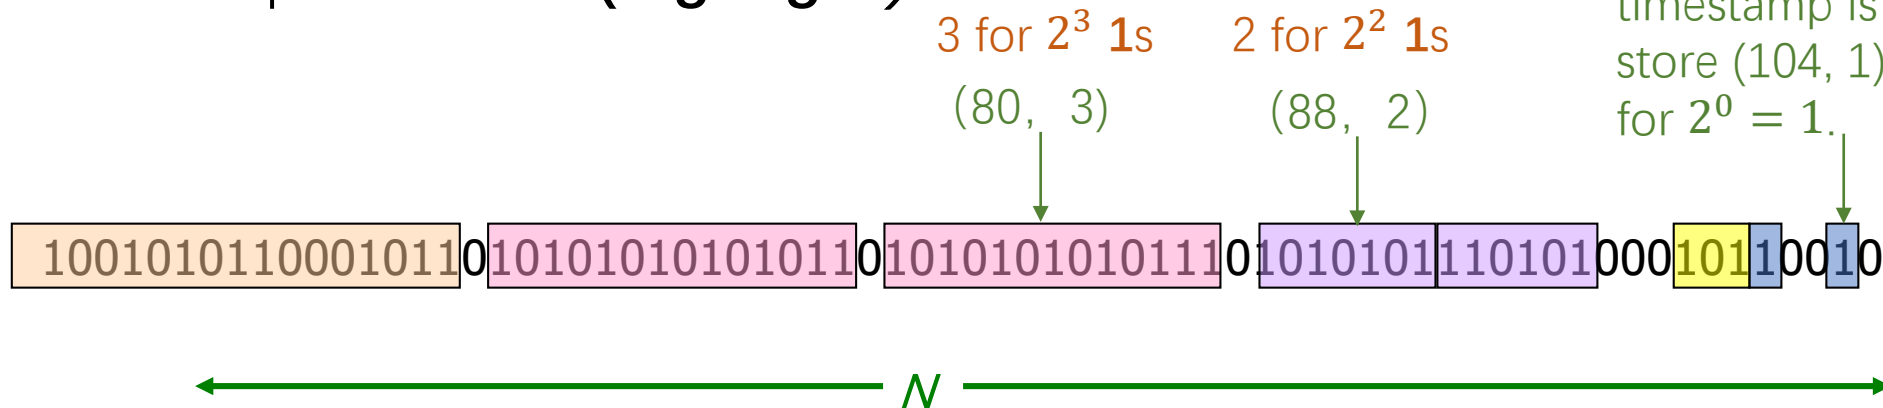
DGIM: Timestamps

- Each bit in the stream has a *timestamp*, starting **1, 2, ...**
- Record timestamps **modulo N** (**the window size**), so we can represent any **relevant** timestamp in **$O(\log_2 N)$** bits




DGIM: Buckets

- A *bucket* in the DGIM method is a record consisting of:
 - The timestamp of its end [$O(\log N)$ bits]
 - The number of 1s between its beginning and end [$O(\log \log N)$ bits]
- **Constraint on buckets:**
Number of **1s** must be a power of 2
 - That explains the $O(\log \log N)$



Representing a Stream by Buckets

- Either **one** or **two** buckets with the same **power-of-2 number** of **1s**
- Buckets **do not overlap** in timestamps
- Buckets are **sorted** by **size**
 - Earlier buckets are not smaller than later buckets
- Buckets **disappear** when their end-time is $> N$ time units in the past



Updating Buckets

- When a new bit comes in, **drop** the last (oldest) bucket if its end-time is **prior to N** time units before the current time
- **2 cases:** Current bit is **0** or **1**
- **If the current bit is 0:**
no other changes are needed

Updating Buckets

- **If the current bit is 1:**
 - **(1)** Create a new bucket of size **1**, for just this bit
 - **End timestamp = current time**
 - **(2)** If there are now three buckets of size 1, combine the oldest two into a bucket of size 2
 - **(3)** If there are now three buckets of size 2, combine the oldest two into a bucket of size 4
 - **(4)** And so on ...

Example: Updating Buckets

Current state of the stream:

100101011000101101010101010101101010101010111010101011101010111010100010110010

Bit of value 1 arrives

001010110001011010101010101011010101010101110101010111010101110101000101100101

Two smallest buckets get merged into a size-2 bucket

00101011000101101010101010101101010101010111010101110101000101100101

Next bit 1 arrives, new size-1 bucket is created, then 0 comes, then 1:

01011000101101010101010101011010101010111010101110101000101100101101

Buckets get merged...


01011000101101010101010101011010101010111010101110101000101100101101

State of the buckets after merging

01011000101101010101010101011010101010111010101110101000101100101101


How to Query?

- To estimate the number of 1s in the most recent N bits:
 1. Sum the sizes of **all** buckets **but the last**
 2. Add **half** the size of the last bucket
- **Remember:** We do **not** know how many **1s** of the last bucket are still within the wanted window




Error Bound: Proof

- **Why is error 50%? Let's prove it!**
- Suppose the last bucket has size 2^r
- Then by assuming 2^{r-1} (i.e., half) of its **1s** are still within the window, we make an error of at most 2^{r-1}
- Since there is at least one bucket of each of the sizes less than 2^r , the true sum is at least
 $1 + 2 + 4 + \dots + 2^{r-1} = 2^r - 1$
- Thus, **relative error at most 50%**



Further Reducing the Error

- Instead of maintaining **1** or **2** of each size bucket, **we allow either $s-1$ or s buckets ($s > 2$)**
 - Except for the largest size buckets, where we can have any number between **1** and **s** of those.



- **Error is at most $\frac{2^{r-1}}{(s-1)(2^r-1)} = O(1/s)$**
- By picking **s** appropriately, we can tradeoff between number of bits we store and the error

Extensions

- Can we handle the case where the stream is not bits, but **integers**, and we want the **sum** of the last k elements?
- We want the sum of the last k elements
 - Amazon: Avg. price of last k sales
- **Solution:**
 - If you know all have at most m bits
 - Treat m bits of each integer as a separate stream
 - Use DGIM to count **1s** in each integer
 - The sum is $= \sum_{i=0}^{m-1} c_i 2^i$

111111110000000011101010101011010101010101110101010111010100010110010
111111110000000011101010101011010101000101110101010111010100011110011

Two streams represent 1